


# ***Advanced Software Engineering***

## **10. Test-Driven Software Experimentation**

**Fall Semester 2024**



Marcus Kessel [marcus.kessel@uni-mannheim.de](mailto:marcus.kessel@uni-mannheim.de), Colin Atkinson [colin.atkinson@uni-mannheim.de](mailto:colin.atkinson@uni-mannheim.de)



# Agenda

- Program Analysis
  - Static vs Dynamic
- Test-Driven Software Experimentation – Software testing is an empirical, continuous activity
- Stimulus Response Matrices (SRMs)
- Differential Testing
  - Test oracles
  - Regression testing
  - (Program-based) Mutation testing
  - N-version testing
- Stimulus Response Hypercubes (SRHs)
- Measuring Code Coverage



# Static vs Dynamic Program Analysis

## Static Analysis

- Analyzes code without execution
  - reasons over abstract code representations (e.g., AST, PDGs)
- No run-time activity required
- Analyzes all possible paths (if reachable)
  - → typically more **complete** and **overapproximates**
- Used for code understanding, static code analyzers (e.g., linting, code metrics ...), symbolic execution, model checking etc.

## Dynamic Analysis (Software Testing)

- Analyzes code behaviour at run-time
  - executes the code on a set of inputs, observing behaviour, performance etc.
- Requires run-time activity
- Only observes actual paths taken
  - → typically more **precise** and **underestimates**
- Testing activities (unit, integration, system etc.)

BUT: Rice's theorem states that *for any non-trivial property  $P$  of the set of partial functions (or programs), it is undecidable to determine whether a given program satisfies this property* (cf. halting problem)


→ Approaches complement each other, a combination of both is often used in practice



# Static vs Dynamic Program Analysis – Example



Program  
(code unit)




**Static** Call Graph  
(all reachable paths)  
→ no information about run-time data

	A	B	C	D
1				
2				
3	1			
4	2	create	?stack	
5	3	push	A1	"hello world"
6	4	peek	A1	
	5	size	A1	
	6	pop	A1	
		size	A1	


*Stimuli S*

execute



missed (not exercised)

**Dynamic** Call Graph  
(paths exercised by a set of tests)  
→ monitors run-time behavior



Note: sometimes some subpaths  
are not reachable (e.g., dead code)

# Software Testing – An Empirical Activity


- **Uncertainty and unpredictability** – software is inherently complex and dynamic (i.e., hard to predict all possible outcomes and scenarios)
  - uncertainty requires empirical investigation through experimentation (i.e., software testing)
- **Exploratory nature** – testing is a process of exploration
  - we attempt to uncover previously unknown issues or defects (cf. Dijkstra)
- **Data-driven decision-making** – testing generates vast amounts of data (e.g., test results, traces, metrics etc.) to make informed decisions about the system's behaviour and quality
  - testers heavily rely on this data to inform their strategies and decisions
- **Experimentation and iteration** – testing involves repeated experimentation with different (or equal) inputs, scenarios or configurations to gather evidence about the software's behaviour
  - iterative process allows testers to refine their understanding of the system and identify potential issues
- **Human involvement** – human-centric activity that requires expertise, experience and creativity
  - next to coverage criteria, testers use their empirical knowledge and understanding of software development, programming languages, and testing techniques to design and execute tests

*“Program testing can be used to show the presence of bugs, but never to show their absence!”*



# Software Testing – Continuous Activity

- Software testing is an empirical, continuous activity ...
  - Based on **observation** and **experimentation**
  - Focused on gathering **evidence** and data
  - Able to adapt to changing requirements or new information
  - Influenced by human judgment and experience
- This perspective emphasizes the importance of continuous learning, iteration, and improvement in software testing
  - aided by automated testing (cf. agile coding practices)
    - unit testing frameworks, continuous integration



Test-Driven Software Experimentation

# Test-Driven Software Experimentation

- A Test-Driven Software Experiment (TDSE) involves the execution of software subjects and the observation and analysis of their “de facto” run-time behaviour (i.e., run-time semantics)
  - *“Test-Driven Software Experiments (TDSEs) are experiments that involve controlled testing of software subjects (i.e., code modules) under various conditions, revealing important properties of the code’s run-time behavior that cannot be predicted solely through static analysis”*
  - Controlled testing: Repeatable observations for (deterministic) software in a controlled environment
    - different environments may cause different observations
- Practitioners
  - Evaluate code and tools, making informed decisions for adoption and integration (e.g., code recommendation)
- Researchers
  - Empirically validate tools and techniques involving the execution of software subjects (e.g., benchmarking tools like test generators)
- Educators
  - Teaching activities (e.g., test-driven exercises)

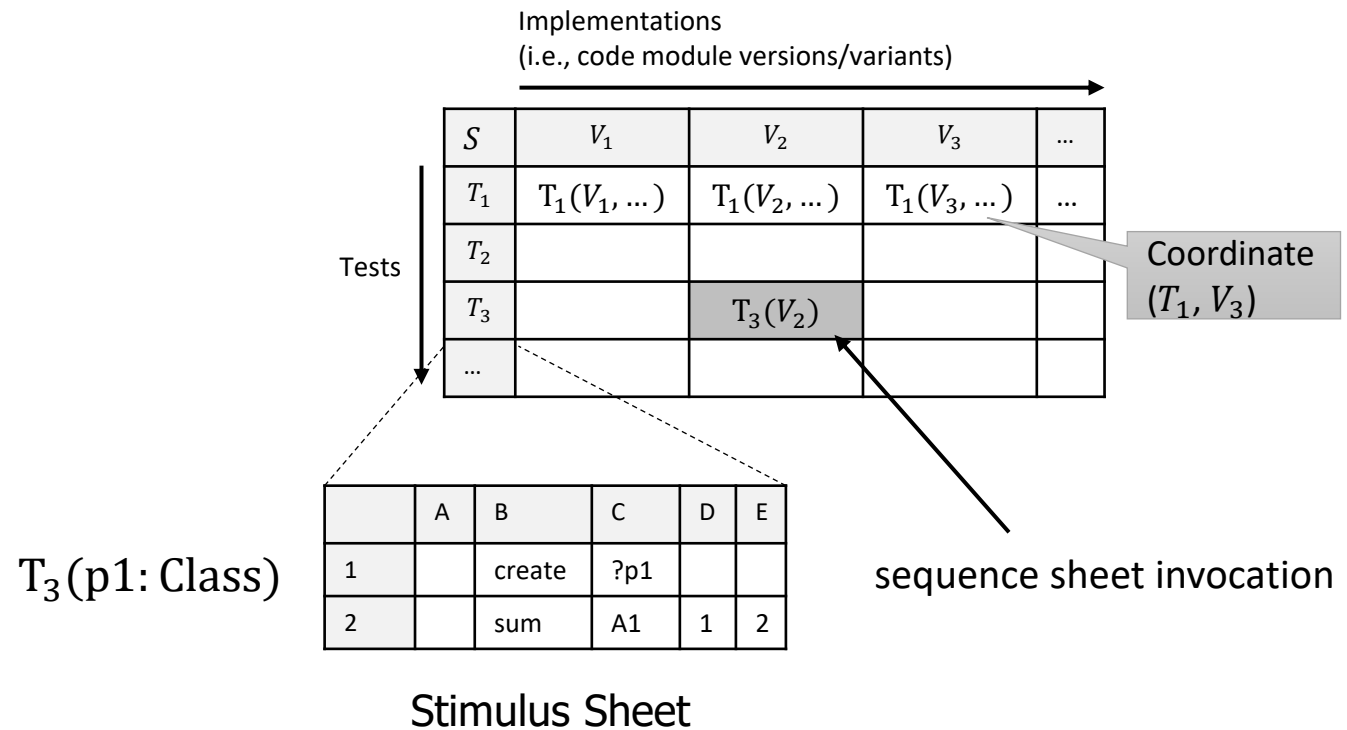


Supporting data structures needed



# Stimulus Matrices (SMs) – Input

- Typically for one functional abstraction  $S$ 
  - i.e., functionality under consideration
    - given by its interface
    - desired behaviour in terms of set of input/output mappings (e.g., sequence sheets)
- Configuration of test/code implementations pairs for execution
  - Columns: One or more Implementations  $V_i$
  - Rows: One or more Tests  $T_j$  (sequence sheets)
    - Typically sequence sheets (Normal, parameterized)

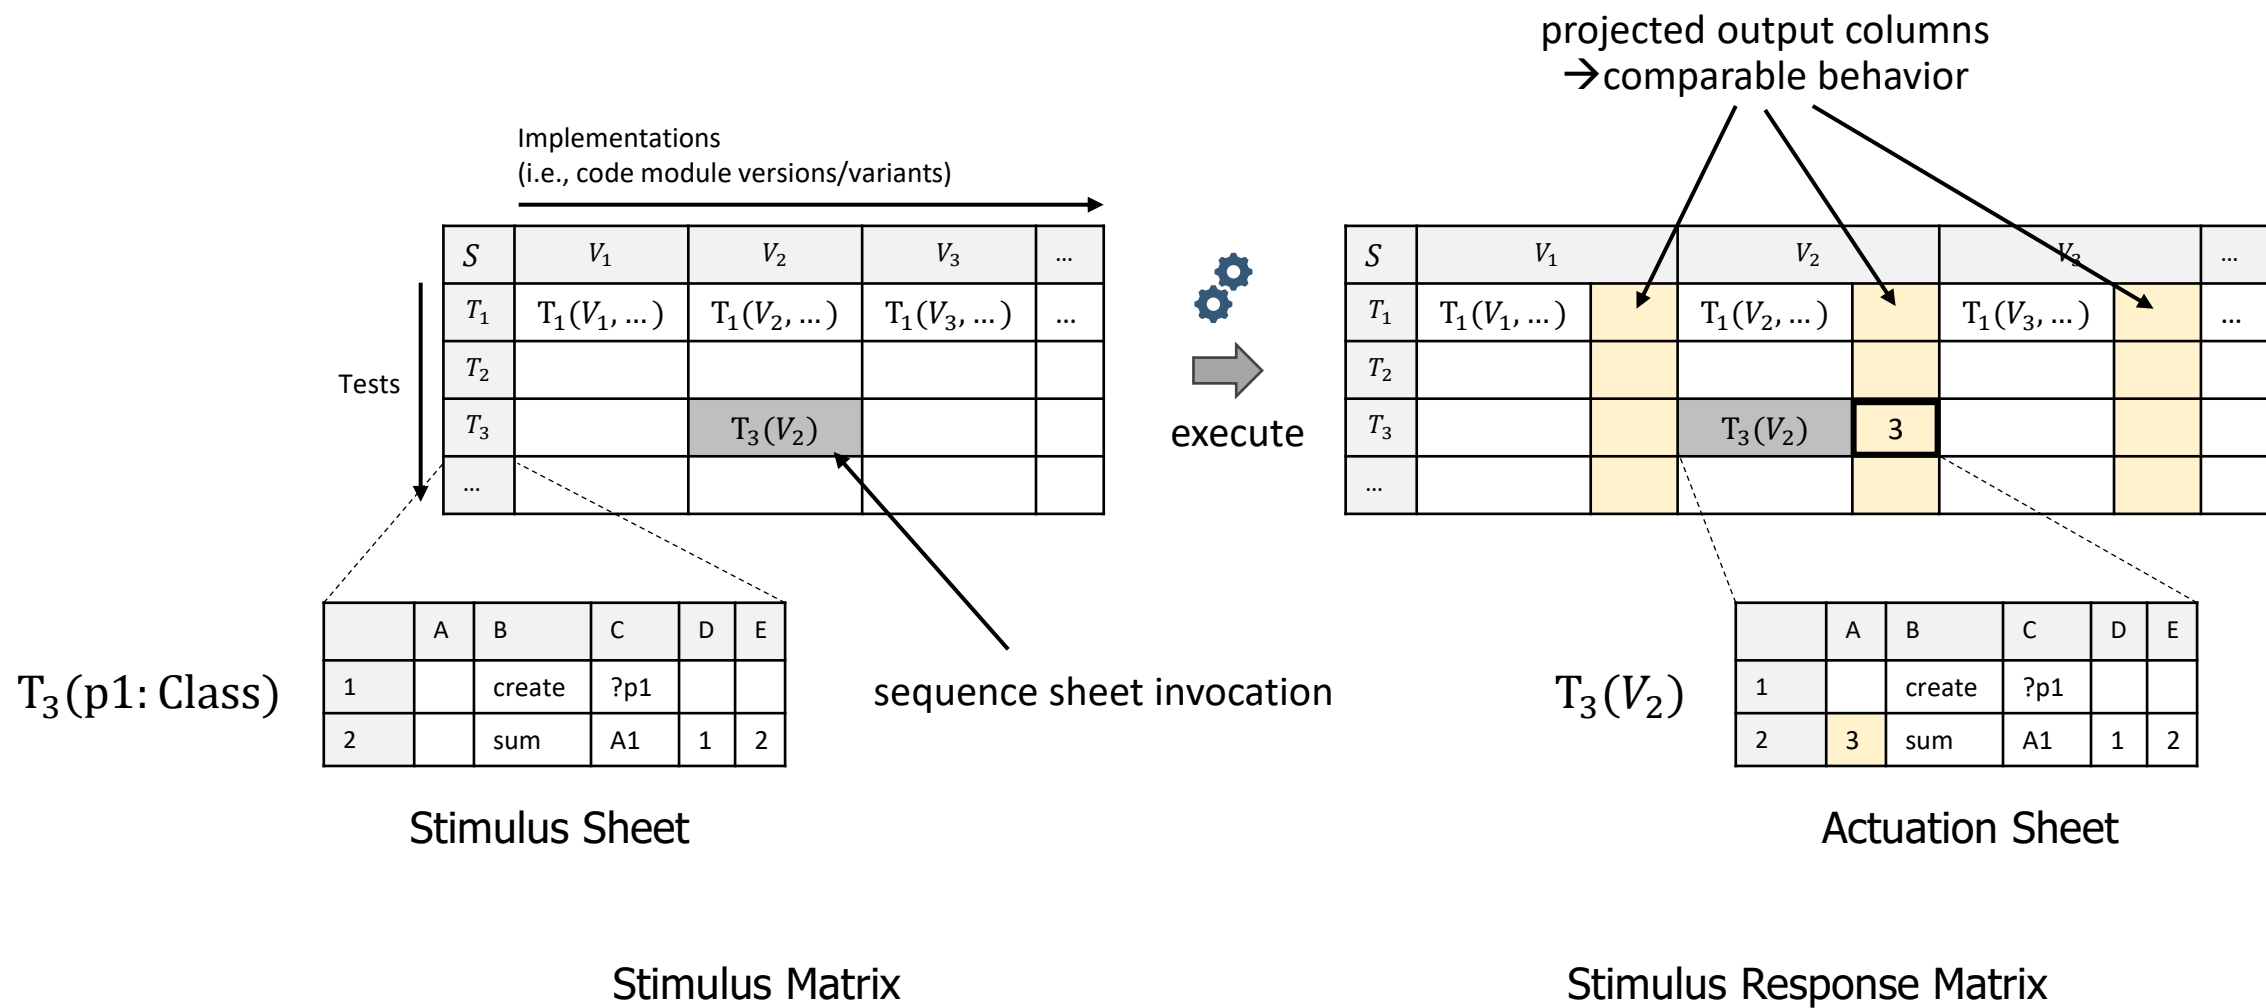




# Stimulus Response Matrices (SRMs) – Output


- Stimulus Response Matrices store run-time responses

Note: Assume that we ignore the first row



# Stimulus Response Matrices (SRMs) – Records

- In general, SRMs store all cells of sequence sheets using two views
  - *black box*: sequence sheet body hidden
  - *white box*: sequence sheet body visible
  - can navigate SRM and select sequence sheet rows/columns/cells of interest
- In addition to functional records (output column), SRMs can also store non-functional information



# SRM Example – Greatest Common Divisor

Problem: GCD (Greatest Common Divisor)  
Interface: `gcd(int, int)->int`

`test(p1:class, p2:int, p3:int)`

	A	B	C	D	E
1		create	?p1		
2		gcd	A1	?p2	?p3

*output    operation    inputs*

Tests

*N versions (code module implementations)*

	V1	V2	V3	...	V <sub>n</sub>
test1	<code>test(V1,03,07)</code>	<code>test(V2, 03,07)</code>	<code>test(V3, 03,07)</code>		<code>test(V<sub>n</sub>, 03,07)</code>
test2	<code>test(V1,10,15)</code>	<code>test(V2,10,15)</code>	<code>test(V3,10,15)</code>		<code>test(V<sub>n</sub>,10,15)</code>
test3	<code>test(V1,49,14)</code>	<code>test(V2,49,14)</code>	<code>test(V3,49,14)</code>		<code>test(V<sub>n</sub>,49,14)</code>
...					
test <sub>m</sub>	<code>test<sub>m</sub>(V1, x, y)</code>	<code>test<sub>m</sub>(V2, x, y)</code>	<code>test<sub>m</sub>(V3, x, y)</code>		<code>test<sub>m</sub>(V<sub>n</sub>, x, y)</code>

 execute

	A	B	C	D	E
1		create	<b>V1</b>		
2	<b>1</b>	gcd	A1	3	7


	V1	V2	V3	...	V <sub>n</sub>
test1	<code>test(V1,03,07)</code> 1	<code>test(V2,03,07)</code> 7	<code>test(V3,03,07)</code> -7		<code>test(V<sub>n</sub>,03,07)</code> ...
test2	<code>test(V1,10,15)</code> 5	<code>test(V2,10,15)</code> 15	<code>test(V3,10,15)</code> -15		<code>test(V<sub>n</sub>,10,15)</code> ...
test3	<code>test(V1,49,14)</code> 7	<code>test(V2,49,14)</code> 14	<code>test(V3,49,14)</code> 28		<code>test(V<sub>n</sub>,49,14)</code> ...
...					
test <sub>m</sub>	<code>test<sub>m</sub>(V1, x, y)</code> ...	<code>test<sub>m</sub>(V2, x, y)</code> ...	<code>test<sub>m</sub>(V3, x, y)</code> ...		<code>test<sub>m</sub>(V<sub>n</sub>, x, y)</code> ...

Note: Assume that we ignore the first row

*comparison of observed outputs*

# Queue Example – Complex Sequences

- Assume the functional abstraction of a queue data abstraction
  - SRMs store a de-collapsed view on rows of sheets (e.g., output column)



	Output	Operation	Inputs (incl. service)	
	A	B	C	D
1	<instance>	create	`MyQueue	
2	TRUE	enqueue	A1	1
3	TRUE	enqueue	A1	2
4	1	peek	A1	
5	2	size	A1	
6	1	dequeue	A1	
7	1	size	A1	

Actuation Sheet

test1(MyQueue)

De-collapsed  
sheet rows

	MyQueue	...	V <sub>n</sub>
test1	test1(MyQueue)	<instance>	...
...2		TRUE	...
...3		TRUE	...
...4		1	...
...5		2	...
...6		1	...
...7		1	...
test2	test2(MyQueue)		...
...	...		...

Stimulus Response Matrix

# Differential Testing & SRMs

- In general, **Differential testing** is a software testing technique that involves comparing the behaviour of two or more versions of a program, typically to detect differences in their functionality
- In practice, many testing activities can be related to differential testing
  - **Test oracles** – compare expected behaviour to actual behaviour of unit under test
    - Specified oracles (input/output mappings, pre-/post conditions, specification lang. ...)
    - Cross-checking oracles – using different implementations of the same functionality
  - **Regression testing** – compare old version of code unit with new version (i.e., verify that code modifications do not change run-time semantics)
  - **Mutation testing** – compare original version of code unit to mutated versions
  - **N-version testing** – compare behaviour of two or more versions of a program and settle (agree) on common behaviour (e.g., critical, dependable software)

 SRMs support and facilitate differential testing

# Test Oracles – Specified Oracle

- A **test oracle** uses expected outputs (“ground truth”) to decide whether a test passed or failed
  - represented as a virtual implementation version in SRMs
  - taken from output column of sheets

	A	B	C	D	E
1		create	V1		
2	1	gcd	A1	3	7

	A	B	C	D	E
1		create	V1		
2	5	gcd	A1	10	15

	A	B	C	D	E
1		create	V1		
2	7	gcd	A1	49	14

Stimulus Sheets

*Expected outputs*

	V1
test1	test(V1,03,07)
test2	test(V1,10,15)
test3	test(V1,49,14)

Stimulus Matrix

Virtual Impl.  
(not executed)



execute

	Oracle		V1	
test1	test(V1,03,07)	1	test(V1,03,07)	1
test2	test(V1,10,15)	5	test(V1,10,15)	5
test3	test(V1,49,14)	7	test(V1,49,14)	11

Stimulus Response Matrix

*Compare expected outputs with actual, observed outputs*

*Actual (observed) outputs*

	A	B	C	D	E
1		create	V1		
2	1	gcd	A1	3	7

	A	B	C	D	E
1		create	V1		
2	5	gcd	A1	10	15

	A	B	C	D	E
1		create	V1		
2	11	gcd	A1	49	14

Actuation Sheets



# Test Oracles – Cross-Checking Oracle

- Reference system *Ref* serves as cross-checking oracle in SRM
  - available reference implementation (e.g., different sort algorithms should return the same outputs for inputs)

	A	B	C	D	E
1		create	<b>V1</b>		
2		gcd	A1	3	7

	A	B	C	D	E
1		create	<b>V1</b>		
2		gcd	A1	10	15

	A	B	C	D	E
1		create	<b>V1</b>		
2		gcd	A1	49	14

Stimulus Sheets

	Ref	V1
test1	<i>test(V1,03,07)</i>	<i>test(V1,03,07)</i>
test2	<i>test(V1,10,15)</i>	<i>test(V1,10,15)</i>
test3	<i>test(V1,49,14)</i>	<i>test(V1,49,14)</i>

Stimulus Matrix

Reference System  execute

	Ref		V1	
test1	<i>test(V1,03,07)</i>	1	<i>test(V1,03,07)</i>	1
test2	<i>test(V1,10,15)</i>	5	<i>test(V1,10,15)</i>	5
test3	<i>test(V1,49,14)</i>	7	<i>test(V1,49,14)</i>	11

Stimulus Response Matrix

*Compare observed outputs*

*Actual (observed) outputs*

	A	B	C	D	E
1		create	<b>V1</b>		
2	1	gcd	A1	3	7

	A	B	C	D	E
1		create	<b>V1</b>		
2	5	gcd	A1	10	15

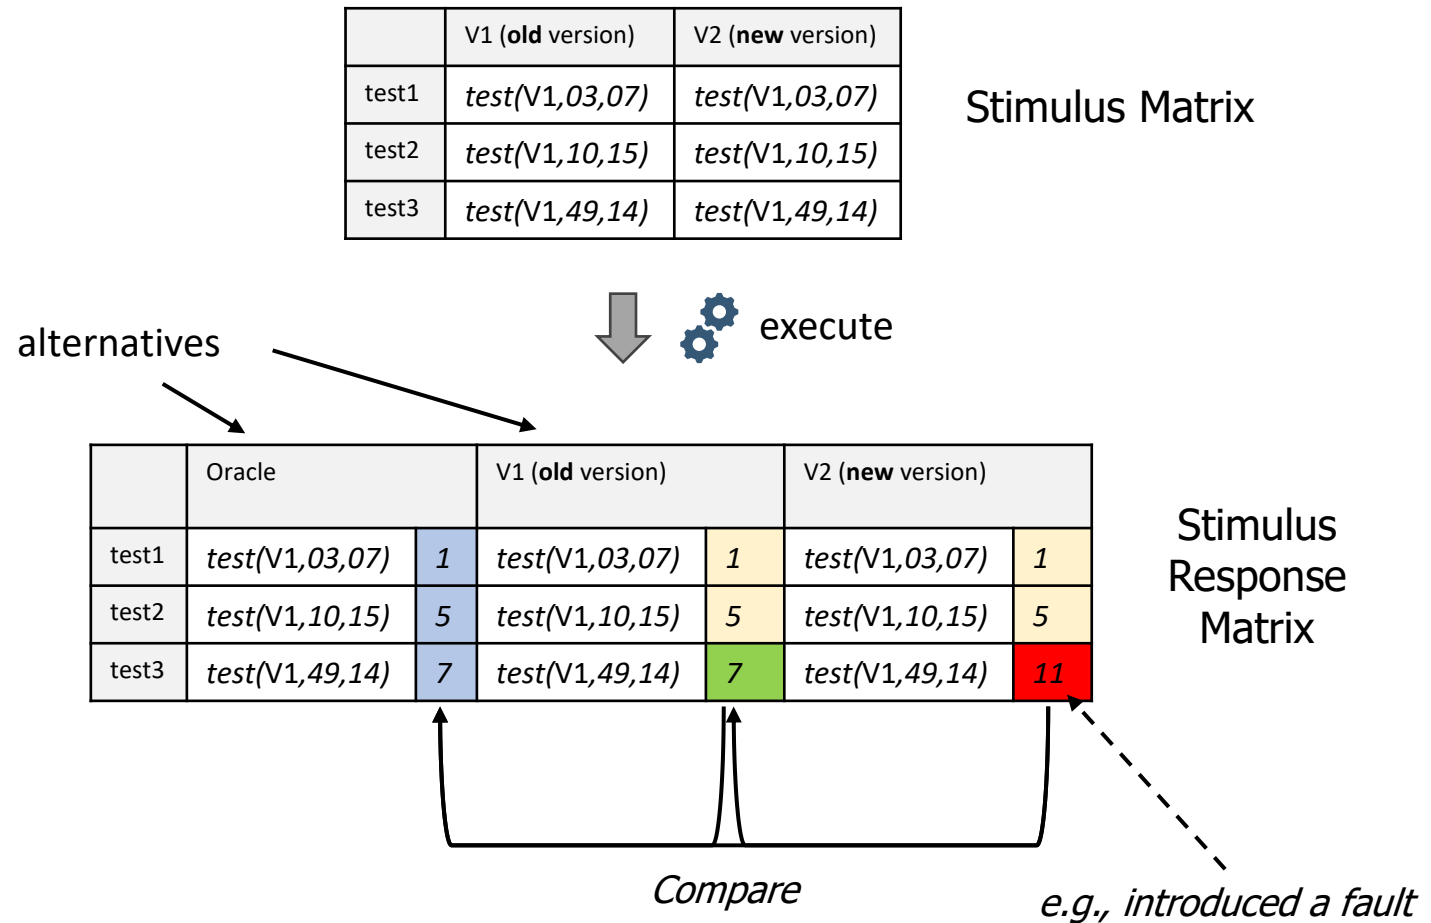
	A	B	C	D	E
1		create	<b>V1</b>		
2	11	gcd	A1	49	14

Actuation Sheets



# Regression Testing

- Regression testing verifies that the software still behaves as expected (i.e., functions correctly) after updates
  - added features
  - fixed faults
  - refactoring, etc.
- Typical realizations
  - re-run existing tests on new version
  - old version may become your (cross-checking) oracle





# Mutation Testing – Program-based Grammars

- Recall syntax-based testing of grammars
- The original and most widely known application of syntax-based testing is to **modify programs**
- **Operators** modify a ground string (program under test) to create **mutant programs**
  - Mimic typical programmer mistakes (e.g., incorrect variable name)
  - Encourage common test heuristics (e.g., cause expressions to be 0)
- Mutant programs must compile correctly (**valid strings**)
- Mutants are **not tests**, but used **to find tests** → assessing test set quality
  - **Fundamental Premise:** In practice, if the software contains a fault, there will usually be a set of mutants that can be killed only by a test case that also detects that fault
- Once mutants are defined, **tests** must be found to cause mutants to fail when executed
- This is called “**killing mutants**”



# Mutation Testing – Operators Example

1. ABS — Absolute Value Insertion
2. AOR — Arithmetic Operator Replacement
3. ROR — Relational Operator Replacement
4. COR — Conditional Operator Replacement
5. SOR — Shift Operator Replacement
6. LOR — Logical Operator Replacement
7. ASR — Assignment Operator Replacement
8. UOI — Unary Operator Insertion
9. UOD — Unary Operator Deletion
10. SVR — Scalar Variable Replacement
11. BSR — Bomb Statement Replacement

## 2. AOR — Arithmetic Operator Replacement:

Each occurrence of one of the arithmetic operators  $+$ ,  $-$ ,  $*$ ,  $/$ , and  $\%$  is replaced by each of the other operators.

Examples:

```
a = m * (o + p);  
 $\Delta 1$  a = m + (o + p);  
 $\Delta 2$  a = m * (o * p);  
 $\Delta 3$  a = m / (o + p);
```

*Remember: Mutated strings (here code) should be close to the original ...*

→ Researchers select appropriate operators through empirical assessments



# Mutation Testing – Killing Mutants

Given a mutant  $m \in M$  for a ground string program  $P$  and a test  $t$ ,  $t$  is said to kill  $m$  if and only if the output of  $t$  on  $P$  is different from the output of  $t$  on  $m$ .

- If mutation operators are designed well, the resulting tests will be very powerful
- Different operators must be defined for different programming languages and different goals
- Testers can keep adding tests until all mutants have been killed
  - Dead mutant: A test has killed it
  - Stillborn mutant: Syntactically illegal
  - Trivial mutant: Almost every test can kill it
  - Equivalent mutant: No test can kill it (same behavior as original)



# Mutation Testing – Test Requirements

Mutation Coverage (MC) : For each  $m \in M$ , TR contains exactly one requirement, to kill  $m$ .

- Recall the conditions from lecture *L2* for a **software failure** to be observed -
  - **Reachability** : The test causes the faulty statement to be reached (in mutation – the mutated statement)
  - **Infection** : The test causes the faulty statement to result in an incorrect state
  - **Propagation** : The incorrect state propagates to incorrect output
  - + **Revealability** (observability) : The tester must observe part of the incorrect output
- Together, these are also known as the RIPR model
- This leads to two variants of mutation coverage ...



# Mutation Testing – Strong vs Weak Mutation Coverage

- Strongly killing mutants – satisfies all conditions (assumed in the remainder ...)

Given a mutant  $m \in M$  for a program  $P$  and a test  $t$ ,  $t$  is said to *strongly kill*  $m$  if and only if the **output** of  $t$  on  $P$  is different from the output of  $t$  on  $m$

- Weakly killing mutants

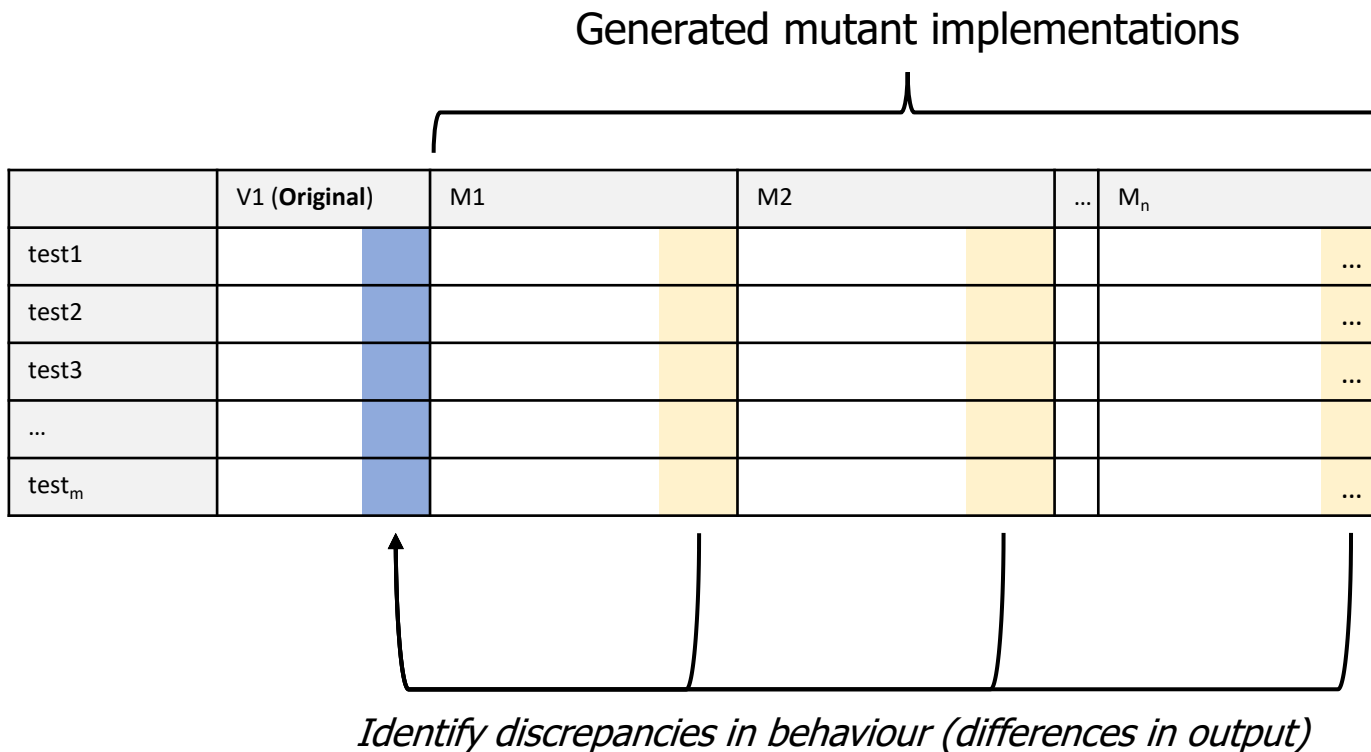
Given a mutant  $m \in M$  that modifies a location  $l$  in a program  $P$ , and a test  $t$ ,  $t$  is said to *weakly kill*  $m$  if and only if the **state** of the execution of  $P$  on  $t$  is different from the state of the execution of  $m$  on  $t$  immediately after  $l$

- Weakly killing satisfies **reachability** and **infection**, but not **propagation**



# Mutation Testing – SRMs

- Basic SRM setting for strong mutation testing



# Mutation Testing – Min Example

## Original Method

```
int min (int A, int B)
{
    int minVal;
    minVal = A;
    if (B < A)
    {
        minVal = B;
    }
    return (minVal);
} // end Min
```

6 mutants

Each represents a  
separate program

## With Embedded Mutants

```
int min (int A, int B)
{
    int minVal;
    minVal = A;
    Δ 1 minVal = B;
    if (B < A)
    Δ 2 if (B > A)
    Δ 3 if (B < minVal)
    {
        minVal = B;
    Δ 4 Bomb ();
    Δ 5 minVal = A;
    Δ 6 minVal = failOnZero (B);
    }
    return (minVal);
} // end Min
```

Replace one variable  
with another

Replaces operator

Immediate runtime  
failure ... if reached

Immediate runtime  
failure if B==0, else  
does nothing



# Mutation Testing – SRM Min Example

- Setting for Min Example
- Assume there are 3 tests, and 6 mutants (see slide before)
  - a mutant is killed if there is at least one test with at least one disagreement in the expected outputs

	A	B	C	D	E
1		create	?c		
2	1	min	A1	2	1

	A	B	C	D	E
1		create	?c		
2	1	min	A1	1	2

	A	B	C	D	E
1		create	?c		
2	1	min	A1	1	1

Generated mutant implementations

	V1 (Original)		M1	M2	M3	M4	M5	M6
test1	test1(V1)	1	1	2	1	E	2	
test2	test2(V1)	1	2	2	1	1	1	
test3	test3(V1)	1	1	1	1	1	1	

➡ Killed  
4 Mutants

Expected outputs ↗ Stimulus Sheets



# Mutation Testing – Equivalent Mutants

- Mutant 3 (M3) in the *min()* example is behaviourally equivalent to the original implementation:

```
minVal = A;  
if (B < A)  
  Δ 3 if (B < minVal)
```

- The infection condition is “(B < A) != (B < minVal)”
- However, the previous statement was “minVal = A”
- Substituting, we get: “(B < A) != (B < A)”
  - This is a logical contradiction !
- Thus no input can kill this mutant

# Mutation Score & SRMs – Example

- We can work out the mutation score directly from SRMs
- Mutation Score
  - Number mutants generated for killing (= 6)
    - i.e., set of test requirements
  - Number of killed mutants (= 4)
  - Score:  $\frac{4}{6} \rightarrow \sim 66\%$
- Important: that we can't reach 100% because of equivalent mutant M3
- Generally, only disagreements to explicit expected values are considered (i.e., marked in blue) – if missing, no comparison is made
  - similar to unit testing assumptions (i.e., missing assertion checks)

	V1 (Original)		M1	M2	M3	M4	M5	M6
test1	test1(V1)	1	1	2	1	E	2	1
test2	test2(V1)	1	2	2	1	1	1	1
test3	test3(V1)	1	1	1	1	1	1	1

one output discrepancy in test2 (column)  
→ M1 killed

two output discrepancies in test1, test2 (column)  
→ M2 killed

# Mutation Testing – Final Words

- Mutation is widely considered the strongest test criterion (“gold standard”)
  - And most expensive ! (run-time requirements in larger code bases ...)
  - By far the most test requirements (each mutant)
  - Usually the most tests
- Mutation subsumes other criteria by including specific mutation operators
- Tool examples
  - Java: *PIT* <https://pitest.org/> , Python: *mutpy*, *mutmut* ...
  - Provide reports (may report other criteria as well) etc. → lookup their suggested default set of mutation operators
  - Can be integrated into automated testing workflows (i.e., continuous integration)
    - Typically compatible to existing unit testing frameworks/drivers

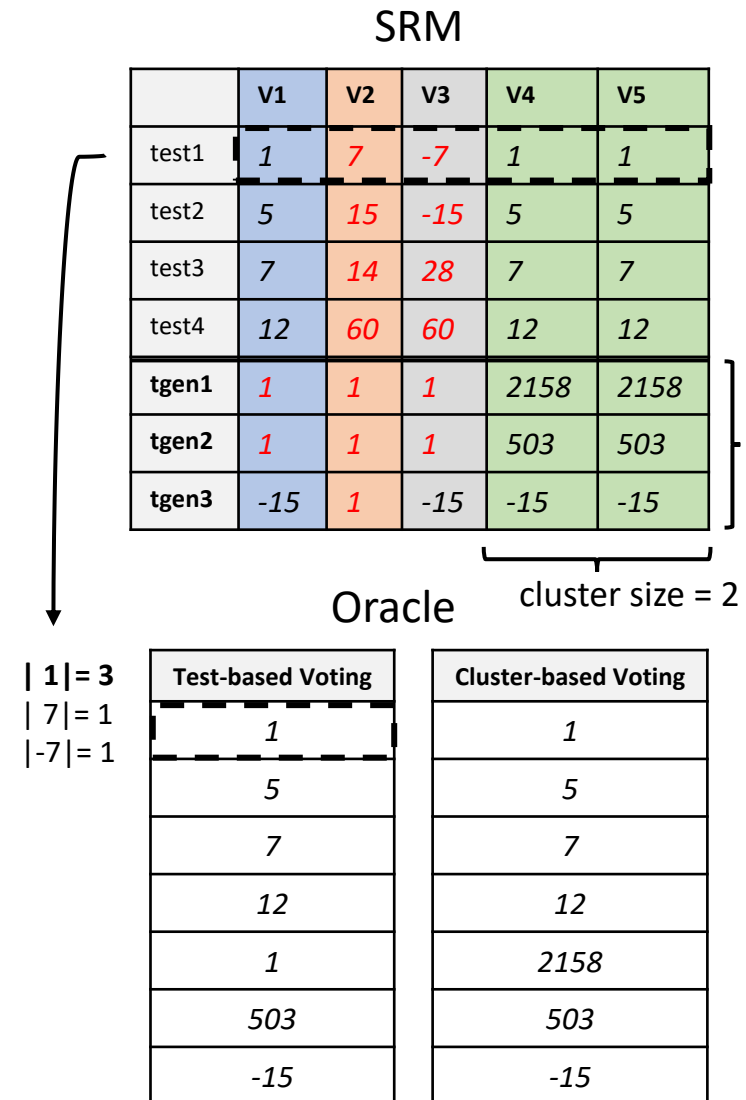
```
122 // Verify for a "." component at next iter
123 3 if ((newcomponents.get(i)).length() > 0 {
124 {
125     newcomponents.remove(i);
126     newcomponents.remove(i);
127 1 i = i - 2;
128 1 if (i < -1)
129 {
130     i = -1;
131 }
132 }
133 }
```

PIT reports

(green line coverage, red mutation coverage)

# N-version Testing

- compare behaviour of two or more versions of a program and settle (agree) on common behaviour (i.e., responses in terms of outputs)
- Traditional hypothesis: multiple versions have been developed by independent developers for the same behavior specification
- assumes that different developers make different mistakes
- Proposed for critical, dependable software (e.g., avionics)
  - correct behavior is identified by majority vote at run-time (for instance)
  - other voting mechanisms possible
- can also be used for practical assessment of multiple implementation candidates where one final candidate is to be selected
  - i.e., fit-for-purpose evaluations in code recommenders
  - however, may require adaptation of the code candidates interfaces




# Software Code Evolution – Continuous Integration

- Test-driven development methods work best when the current version of the software can be run against all tests at any time

A *continuous integration server* rebuilds the system, returns, and reverifies tests whenever *any* update is checked into the repository

- Mistakes are caught earlier
- Other developers are aware of changes early
- The rebuild and reverify must happen as soon as possible
- Thus, tests need to execute quickly

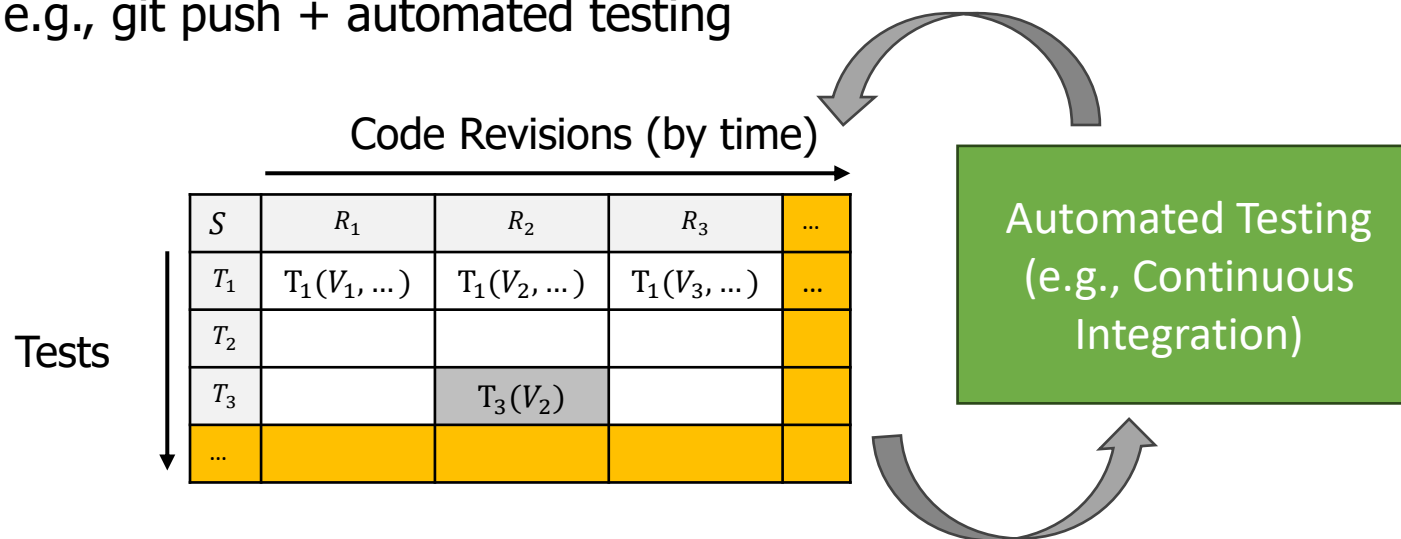


e.g., Jenkins weather report

A *continuous integration server* doesn't just run tests, it decides if a modified system is still correct

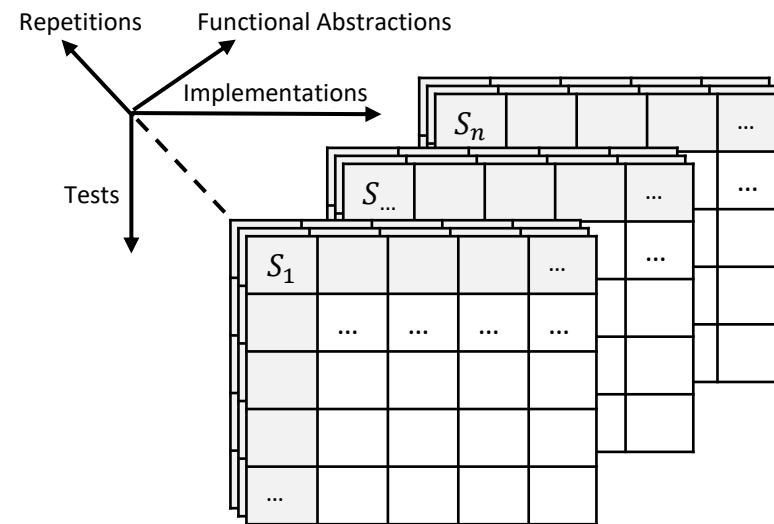
# Software Code Evolution – Continuous Lifecycle

- Over time, as the production code evolves, the set of tests for the code implementation evolves as well
  - tests may be added (e.g., “bug-fixing” tests), updated or removed to maintain/improve reliability
- if changes are made to the implementation
  - code revisions (new versions) are added from time to time and are automatically verified for correctness
    - e.g., git push + automated testing



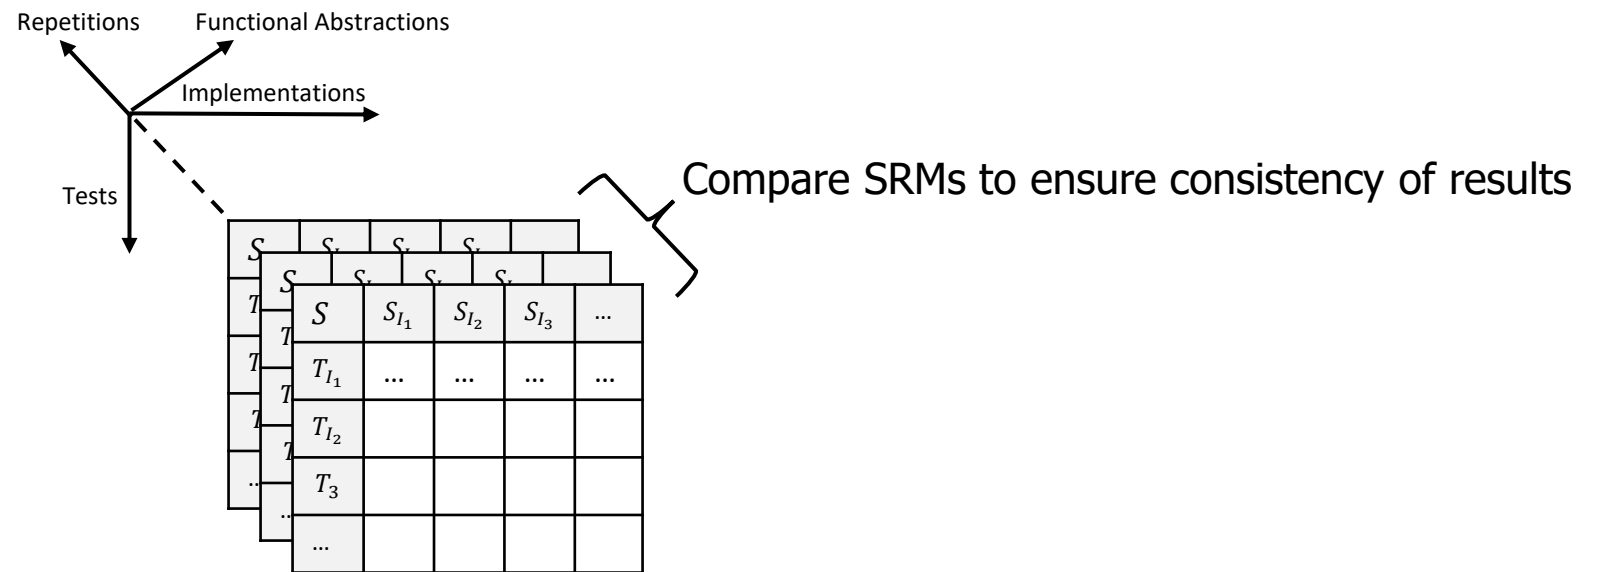
# Stimulus Response Hypercubes – SRHs

- Aggregate multiple SRMs into a cube structure (cf. data warehousing applications)
- Navigate across multiple dimensions to inform decision-making and test reporting
- Functional abstractions  $S_i$ 
  - Tests
  - Implementations
- Repetitions
- ...
- Enables software analytics
  - Reason over data (e.g., metrics etc.)



# Repeatability – Checking for Test Flakiness

- Test flakiness refers to the phenomenon where automated tests occasionally fail or behave unexpectedly, even when there are no actual issues with the code being tested
  - Random failures
  - Inconsistent behaviour
  - False positives/negatives
- Causes, for instance
  - Source of randomness
  - Concurrency issues
  - External services (e.g., DBs)
- Consequences, for instance
  - Maintenance costs
  - Confidence erodes
  - May mask actual issues with the code





# Advanced Topics – An Example

- Test set minimization of a test set T
  - Problem: Imagine you code base as well as the test sets evolve over time → Scalability?
  - Goal: Identify a minimal set of tests to improve efficiency, thus minimize run-time costs (time, money etc.) – also see *test set prioritization* (only running those tests affected by a change)
  - Idea: if a pair of tests satisfy the same test requirement, we can drop either of these, while still satisfy the same test requirements
    - E.g., identifying redundant tests based on the killing of the same mutants for a pair of tests → drop one redundant test
  - Similar to mutation matrices



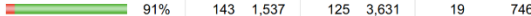





















	V1 (Original)	M1	M2	M3	M4
test1		1	1	7	1
test2		5	0	5	5
test3		7	7	1	7

assume that test1 and test2 kill the same mutant M3 → drop test1 or test3

# Measuring Code Coverage in Practice

- a way to quantify how much of a program's code is exercised during automated tests at run-time
- provides insight into the thoroughness of the test suite by measuring the percentage of code that has been exercised
- can be used to identify non-exercised code elements
- Typical types of code coverage
  - Line coverage (or statement coverage)
  - Branch coverage
  - Method/function or class coverage
- Typical metrics reported (note: should be carefully interpreted!)
  - Percentage of code covered/missed etc.
- Tool examples
  - Java: *JaCoCo*, (Open)*Clover*, ...
  - Python: *Coverage.py*, ...

## JaCoCo

Element	Missed Instructions	Cov.	Missed Branches	Cov.	Missed Cxty	Missed Lines	Missed Methods	Missed Classes
 org.jacoco.core		97%		91%	143 1,537	125 3,631	19 746	2 147
 org.jacoco.examples		58%		64%	24 53	97 193	19 38	6 12
 org.jacoco.agent.rt		75%		83%	32 130	75 344	21 80	7 22
 jacoco-maven-plugin		90%		82%	35 193	49 465	8 116	1 23
 org.jacoco.cli		97%		100%	4 109	10 275	4 74	0 20
 org.jacoco.report		99%		99%	4 572	2 1,345	1 371	0 64
 org.jacoco.ant		98%		99%	4 162	8 428	3 110	0 19
 org.jacoco.agent		86%		75%	2 10	3 27	0 6	0 1
Total	1,438 of 28,925	95%	183 of 2,386	92%	248 2,766	369 6,708	75 1,541	16 308



# Conclusion

- Software testing is an empirical activity in general → Test-Driven Software Experimentation
  - Testing is done in a continuous manner (tests co-evolve with the units under test)
- Differential testing with SRMs
  - Test oracles (specified and cross-checking oracles)
  - Regression testing
  - Mutation testing / score
  - N-version testing
- SRHs for continuous testing
  - Repeatability and test flakiness
- Advanced topics – e.g., test set minimization
- Measuring Code Coverage



# Literature

- Software Testing

- Ammann, Paul, and Jeff Offutt. *Introduction to software testing*. Cambridge University Press, 2017.
- Barr, Earl T., et al. "The oracle problem in software testing: A survey." *IEEE transactions on software engineering* 41.5 (2014): 507-525.
- Chen, Liming, and Algirdas Avizienis. "N-version programming: A fault-tolerance approach to reliability of software operation." *Proc. 8th IEEE Int. Symp. on Fault-Tolerant Computing (FTCS-8)*. Vol. 1. 1978.

- Data Structures

- Publications

- Marcus Kessel, Colin Atkinson, Promoting open science in test-driven software experiments, *Journal of Systems and Software*, Volume 212, 2024, 111971, ISSN 0164-1212, <https://doi.org/10.1016/j.jss.2024.111971>
- Marcus Kessel, LASSO – an observatorium for the dynamic selection, analysis and comparison of software, Dissertation, 2023 - <https://madoc.bib.uni-mannheim.de/64107/>

- LASSO Platform

- <https://softwareobservatorium.github.io/>

